

ECE 454 Course Notes

Distributed Computing

Parsa Torabian

Forked from Michael Socha

University of Waterloo

Spring 2019

Contents

1	Course Overview	1
2	Introduction	2
2.1	What is a Distributed System?	2
2.2	Rationale Behind Distributed Systems	2
2.3	Middleware	2
2.4	Goals of Distributed Systems	2
2.4.1	Supporting Resource Sharing	2
2.4.2	Making Distribution Transparent	3
2.4.3	Being Open	3
2.4.4	Being Scalable	3
2.5	Common Fallacies of Distributed Computing	4
2.6	Types of Distributed Systems	4
3	Architectures	6
3.1	Definitions	6
3.2	Architectural Styles	6
3.2.1	Layered	6
3.2.2	Object-Based Architecture	6
3.2.3	Data-Centered Architecture	6
3.2.4	Event-Based Architecture	7
3.2.5	P2P Systems	7
3.2.6	Hybrid Architectures	7
3.3	Self-Management	7
4	Processes	8
4.1	Controlling Processes in Linux	8
4.2	Context Switching During IPC	8
4.3	Threads	8
4.4	Hardware and Software Interfaces	8
4.5	Interfaces in Network Systems	9
4.6	Server Clusters	9
5	Communication	10
5.1	Access Transparency	10
5.2	Remote Procedure Calls (RPCs)	10
5.2.1	Representation of Values	10
5.2.2	Synchronous vs Asynchronous	11
5.3	Message Queuing Model	11
5.4	Coupling Between Communicating Processes	11
5.4.1	Referential Coupling	11
5.4.2	Temporal Coupling	11
6	Apache Thrift	12

6.1	Overview	12
6.2	Thrift Software Stack	12
6.3	Distribution Transparency	12
6.4	Application Protocol Versioning	12
6.5	IDL Syntax	13
6.6	Todo: Java Server implementations	13
6.7	Some Programming Tips	13
7	Distributed File Systems	14
7.1	Modes of Access	14
7.2	Network File System (NFS)	14
7.3	Google File System (GFS)	14
7.4	File Sharing Semantics	15
8	Hadoop MapReduce	16
8.1	Background	16
8.2	High-level Architecture	16
8.3	MapReduce Basics	16
8.3.1	Mappers	16
8.3.2	Reducers	16
8.4	Technical Definitions	17
8.5	Fault Tolerance	17
8.6	Common Programming Patterns	17
8.6.1	Counting and Summing	17
8.6.2	Selection	18
8.6.3	Projection	18
8.6.4	Index Inversion	18
8.6.5	?Cross-Correlation	18
9	Apache Spark	19
9.1	Background	19
9.2	Lineage	19
9.3	Transformations and Actions	19
9.3.1	Narrow vs Wide Dependencies	19
9.3.2	Pagerank	19
9.4	Spark vs Hadoop MapReduce	19
10	Distributed Graph Processing	21
10.1	Google’s Solution: Pregel	21
10.1.1	Supersteps	21
10.1.2	Initialization	21
10.1.3	Combiners and Aggregators	22
10.1.4	Fault Tolerance	22
10.2	Open-source Implementations	22

11	Consistency and Replication	23
11.1	Replication Overview	23
11.1.1	Purposes of Replication	23
11.1.2	Replicated Data Stores	23
11.2	Consistency Models	23
11.2.1	Sequential Consistency	23
11.2.2	Causal Consistency	23
11.2.3	Linearizability	24
11.2.4	Eventual Consistency	24
11.3	Replication Protocols	24
11.3.1	Primary-Based Replication Protocols	24
11.3.2	Quorum-Based Replication Protocols	25
11.3.3	Eventually-Consistent Replication Protocols	25
12	Fault Tolerance	26
12.1	Dependability	26
12.2	Failures	26
12.3	Errors and Faults	26
12.4	Masking Failure by Redundancy	27
12.5	Consensus Problem	27
12.5.1	Variations	27
12.6	RPC Semantics Under Failures	27
12.6.1	Dealing with RPC Server Crashes	28
13	Apache ZooKeeper	29
13.1	Data Model	29
13.2	Node Flags	29
13.3	Consistency Model	29
13.4	Servers	29
14	Distributed Commits and Checkpoints	30
14.1	Transactions	30
14.2	Two-Phase Commits	30
14.2.1	Recovery from Failures	30
14.2.2	Distributed Checkpoints	30
15	RAFT	32
16	Apache Kafka	33
16.1	Topics and Logs	33
16.2	Producers and Consumers	33
16.3	Stream Semantics	33
16.4	Windowed Streams	34
17	Clock Synchronization	35
17.1	Background on Clocks	35

17.1.1	Calendars	35
17.1.2	Timekeeping Standards	35
17.1.3	Limitations of Atomic Clocks	35
17.2	Terminology	36
17.3	Network Time Protocol (NTP)	36
17.3.1	Clock Strata	36
17.3.2	Precision	36
17.4	Logical Clocks	36
17.4.1	Lamport Timestamping	37
17.4.2	Vector Clocks	37
18	CAP Principle	38
18.1	CP vs AP Systems	38
18.2	PACELC	38
18.3	Quorum-Based Consistency and CAP	38

1 Course Overview

This course is intended to provide an introduction into distributed systems.

Topics covered include:

- Some of the general architectures, protocols, and algorithms underlying modern distributed systems
- Techniques for making distributed systems scalable and dependable, as well as the relevant trade-offs
- Remote procedure call (RPC) frameworks for implementing distributed services
- Scalable data processing frameworks for solving analytics problems
- Fault-tolerant coordination services for configuration management, synchronization, and failure-detection in distributed software systems
- Stream-processing engines for real-time analytics
- Relative merits of distributed versus centralized systems

Open-source software covered:

- Memcached (client-server paradigm)
- Apache Thrift (Remote Procedure Calls)
- Apache Hadoop (Big Data Processing)
- Apache Spark (Big Data Processing)
- Apache Zookeeper (Distributed Coordination)
- Apache Kafka (Data Streaming)

2 Introduction

2.1 What is a Distributed System?

A distributed system is a collection of autonomous computing elements that appear to its users to be a single coherent system.

2.2 Rationale Behind Distributed Systems

Distributed systems are often used for the following reasons:

- Resource sharing
- Integrating multiple systems
- Centralized systems may not be as effective as many smaller systems working together
- Users themselves may be mobile and distributed physically from one another

2.3 Middleware

Middleware is a layer of software that separates applications from their underlying platform. Middleware is often used by distributed systems to support heterogeneous computers and networks while offering a single-system view. Common middleware services include:

- Communication (e.g. add job to remote queue)
- Transactions (e.g. access multiple independent services atomically)
- Service composition (e.g. Independent map system enhanced with independent weather forecast system)
- Reliability (e.g. replicated state machine)

2.4 Goals of Distributed Systems

2.4.1 Supporting Resource Sharing

Resources can include:

- Peripheral devices (e.g. printers, video cameras)
- Storage facilities (e.g. file servers)
- Enterprise data (e.g. payroll)
- Web page (e.g. Web search)

- CPUs (e.g. supercomputers)

2.4.2 Making Distribution Transparent

Distribution transparency attempts to hide that processes and resources are physically distributed. Types of transparency include:

- **Access:** Hide differences in data representation and how data is accessed
- **Location:** Hide where resource is located
- **Migration:** Hide that a resource may move to another location
- **Relocation:** Hide that a resource may move to another location while in use
- **Replication:** Hide that a resource is replicated
- **Concurrency:** Hide that a resource may be shared by users competing for that resource
- **Failure:** Hide the failure and recovery of a resource

2.4.3 Being Open

An open distributed system offers components that can be easily used by or integrated into other systems. Key properties of openness include:

- Interoperability
- Composability
- Extensibility
- Separation of policy from mechanism

2.4.4 Being Scalable

Scalability covers a system's ability to expand along three axes:

- Size (e.g. adding users and resources)
- Geography (e.g. users across large distances)
- Administration (e.g. multiple independent admins)

Centralized systems tend to have limited scalability.

A few common scaling techniques include:

- Hiding communication latencies (i.e. trying to avoid waiting for responses from remote-server machines)

- Partitioning (i.e. taking a component, splitting it into smaller parts, and spreading those parts across the system)
- Replication (i.e. adding multiple copies of a component to increase availability, balance load, support caching, etc.)

2.5 Common Fallacies of Distributed Computing

The following are common beginner assumptions that can lead to major trouble:

- The network is reliable
- The network is secure
- The network is homogeneous
- The network topology is static
- Latency is 0
- Bandwidth is unlimited
- Transport cost is 0
- There is one administrator

2.6 Types of Distributed Systems

- **Web sites and Web services**
- ?Difference from cluster computing? **High performance computing (HPC):** High-performance computing in distributed memory settings, where message communication is used instead of shared memory.
- **Cluster Computing:** Distributing CPU or I/O intensive tasks over multiple servers. Eg: Hadoop
- **Cloud and Grid Computing:** Grid computing focuses on combining resources from different institutions, while cloud computing provides access to shared pools of configurable system resources. Eg: IAAS, PAAS, SAAS
- **Transaction Processing:** Distributed transactions coordinated by transaction processing (TP) monitor.
- **Enterprise Application Integration (EAI):** Middleware often used as communication facilitator in such systems. Eg: Client Apps to EAI Middleware to Enterprise Apps

- **Sensor Networks:** Each sensor in a network can process and store data, which can be queried by some operator. Sensors often rely on in-network data processing to reduce communication costs. Eg: Irrigation

3 Architectures

3.1 Definitions

- **Component:** A modular unit with a well-defined interface
- **Connector:** Mechanism that mediates communication, coordination, or cooperation among components
- **Software architecture:** Organization of software components
- **System architecture:** Instantiation of software architecture in which software components are placed on real machines
- **Autonomic System:** System that adapts to its environment by monitoring its own behavior and reacting accordingly

3.2 Architectural Styles

3.2.1 Layered

Control flows from layer-to-layer; requests flow down the hierarchy and responses flow up, with each layer only interacting with its two neighboring layers. Layered architectures are often used to support client-server interactions, where a client component requests a service from a server component and waits for a response. Many enterprise systems are organized into three layers, namely a user interface, application layer, and database. Note that middle layers may act as a client to the layer below and a server to the layer above.

Vertical distribution describes the logical layers of a system being organized as separate physical tiers. Horizontal distribution describes a single logical layer being split across multiple machines. ?How do these two types of distribution affect performance, scalability, and dependability?

3.2.2 Object-Based Architecture

Components are much more loosely organized than in a layered architecture, with components being able to interact freely with one another and no strict concept of layer.

3.2.3 Data-Centered Architecture

Components communicate by using a shared data repository, such as a database or file system.

3.2.4 Event-Based Architecture

Components communicate by sharing events. Publish/subscribe systems can be used for sharing news, balancing workloads, asynchronous workflows, etc.

In practice, many systems are hybrids of the architectures listed above.

3.2.5 P2P Systems

Peer-to-peer systems rely on horizontal distribution and are designed to deal with churn (machines joining and leaving). They organize processes in a overlay network that defines a set of communication channels. Eg: Chord is a P2P distributed hash table (DHT) that uses a ring overlay?

3.2.6 Hybrid Architectures

BitTorrent combines client-server and P2P architectures. Client nodes obtain tracker information from a server and then exchange data with peer nodes.

3.3 Self-Management

Self-managing systems can be constructed using a feedback loop that monitors system behaviors and adjusts the system's internal operation.

4 Processes

4.1 Controlling Processes in Linux

Below are some useful process-related commands:

- `ps` - lists running processes
- `top` - lists processes with top resource usage
- `kill` / `pkill` / `killall` - used to terminate processes
- `jobs` - lists currently running jobs
- `bg` - backgrounds a job
- `fg` - foregrounds a job
- `nice` / `renice` - sets priority of processes
- `CTRL-C` - stops job in terminal
- `CTRL-Z` - suspends job in terminal (use `fg` or `bg` to resume)

4.2 Context Switching During IPC

Inter-process communication can be used to help coordinate processes. However, IPC can be costly, since it requires a context switch from user to kernel space and back.

4.3 Threads

Threads running in the same process can communicate through shared memory. Threads are sometimes referred to as lightweight processes (LWPs). Multithreading applications often follow a dispatcher/worker design, where a dispatcher thread receives requests and feeds them to a pool of worker threads.

4.4 Hardware and Software Interfaces

Processes and threads may interact with underlying hardware either directly through processor instructions or indirectly through library functions or operating system calls. The general layers of interaction (from more abstract to less) are applications, libraries, operating system, and hardware. These layers may interact with any layer below. Distributed systems often run in virtual environments that manage interactions with lower layers. A major benefit of such virtualization is improved portability. Virtual machine monitors can offer additional benefits, including server consolidation, live migration of VMs to support load balancing and proactive maintenance, and VM replication. ?Drawbacks of VMs?

4.5 Interfaces in Network Systems

Networked applications communicate by exchanging messages. The format of these messages is determined by a message-passing protocol.

4.6 Server Clusters

Servers in a cluster are often organized in three physical tiers. The first is a load balancer, followed by application servers followed by a database or file system.

5 Communication

5.1 Access Transparency

Middleware can be used to provide access transparency for a distributed system, meaning that differences in data representation and how data is accessed can be hidden. This is typically done by middleware isolating the application layer from the transport layer.

5.2 Remote Procedure Calls (RPCs)

RPCs serve as the equivalent of conventional procedure calls, but for distributed systems. RPCs are implemented using a client-server protocol, where a client interacts with a network using a piece of software known as a client stub, and a server interacts with a network using a server stub. The steps of a typically RPC are detailed below:

1. Client process invokes client stub using ordinary procedure call.
2. Client stub builds message, passes it to client OS.
3. Client OS sends message to server OS.
4. Server OS delivers message to server stub.
5. Server stub unpacks parameters, invokes appropriate handler in server process.
6. Handler processes message, returns result to server stub.
7. Server stub packs result into message, passes it to server OS.
8. Server OS sends message to client OS.
9. Client OS delivers message to client stub.
10. Client stub unpacks result, delivers it to client process.

5.2.1 Representation of Values

Packing parameters into a message is known as parameters marshalling (with unpacking known as demarshalling). Although data representations can be different across multiple machines (e.g. big-endian vs little-endian representation), the purpose of marshalling is to represent data in a machine-independent and network-independent format that all communicating parties expect to receive. The signatures of RPC calls can be defined using what is known as an interface definition language (IDL).

5.2.2 Synchronous vs Asynchronous

In a synchronous RPC, the client waits for a return value from the server before resuming execution. In an asynchronous RPC, the client may resume execution as soon as the server acknowledges receipt of the request (client does not need to wait for a return value). A variation of an asynchronous request where a client does not wait to receive any acknowledgment from the server is known as a one-way RPC.

5.3 Message Queuing Model

As an alternative to RPCs, components in a distributed system may also communicate using a message queue that persists messages until they are consumed by a receiver. This technique allows for persistent communication that does not need to be tightly coupled in time. The primitive actions of a simple message queue are:

- **Put:** Append message to queue.
- **Get:** Block until queue not empty, then remove first message.
- **Poll:** Check specified queue for messages, remove the first message. Never block.
- **Notify:** Install handler to be called when message put into queue.

A key disadvantage of message queuing is that the delivery of the message ultimately rests with the receiver, and often cannot be guaranteed. Message queuing follows a design similar to publish-subscribe architectures, and is an example of message-oriented middleware (MOM).

5.4 Coupling Between Communicating Processes

5.4.1 Referential Coupling

Referential coupling means that one process has to explicitly reference another one for them to communicate (e.g. connecting to web server using IP address and port number).

5.4.2 Temporal Coupling

Temporal coupling means that processes must both be running for them to communicate (e.g. client cannot execute RPC if server is down).

6 Apache Thrift

6.1 Overview

Apache Thrift is an IDL first developed by Facebook, and now managed as an open-source project in the Apache Software Foundation. Thrift provides a software stack and code generation engine to support RPCs between applications written in a wide variety of common languages, including C++, Java, Python, and PHP.

6.2 Thrift Software Stack

The elements in the Thrift software stack, from top to bottom, are as follows:

- **Server:** Sets up the lower layers of the stack, and then awaits incoming connections, which it hands off to the processor. Servers can be single or multi-threaded.
- **Processor:** Handles reading and writing IO streams, and is generated by the Thrift compiler.
- **Protocol:** Defines mechanism to map in-memory data structures to wire format (e.g. JSON, XML, compact binary)
- **Transport:** Handles reading to and writing from network (e.g. using HTTP, raw TCP, etc)

6.3 Distribution Transparency

If Thrift clients must know the hostname and port for a given service, location transparency is violated. Also, although IDLs seek to provide access transparency, that too may be violated in certain conditions (e.g. when certain protocol or transport exceptions are thrown).

6.4 Application Protocol Versioning

Thrift fields can be modified and remain compatible with old versions, provided that the rules below are followed:

- Fields are associated with tag numbers, which should never be changed.
- New fields must be optional and provide default values.
- Fields no longer needed can be removed, but their tag numbers cannot be reused.

6.5 IDL Syntax

- **Base Types:** bool, byte, i16, i32, i64, double, binary, string, void
- **Containers:** list, set, map (of types from above)
- **Other types:** const, typedef, enum, struct, exception
- **Field modifiers:** required, optional, default
- **Services and procedures:** service, extends, oneway
- **namespace:** determines java package
- **File inclusion:** Similar to C/C++ preprocessor

6.6 Todo: Java Server implementations

- **TSimpleServer:** uses a single thread and blocking I/O
- **TNonblockingServer:** uses a single thread and non-blocking I/O. It can handle parallel connections but executes requests serially just like TSimpleServer
- **THsHaServer:** uses one thread for network I/O and a pool of worker threads. It can process multiple requests in parallel
- **TThreadPoolServer:** uses one thread to accept connections and then handles each connection using a dedicated thread drawn from a pool of worker threads.
- **TThreadedSelectorServer:** uses a pool of threads for network I/O and a pool of worker threads for request processing.

6.7 Some Programming Tips

- To separate policy from mechanism, it is generally a bad idea to hardcode hostnames and port numbers; it is usually preferable to accept command line arguments or use a property file.
- Objects built on top of TCP/IP connections (e.g. TSocket, TServerSocket) should be reused if possible to avoid overhead of establishing and tearing down connections.
- By default Thrift transports, protocols and client stubs are not thread safe; different threads should share these items carefully, or not share them at all.

7 Distributed File Systems

7.1 Modes of Access

- **Remote access model:** Client sends requests to access file stored on remote server.
- **Upload/download model:** File is downloaded from remote server, processed on client, and uploaded back to the server.

7.2 Network File System (NFS)

NFS is a DFS first developed by Sun Microsystems in 1984, and remains heavily used in Unix-like systems. Features include:

- Client-side caching to reduce client-server communication
- Delegation of files from a server to a client, where a client can temporarily store and access a file, after which the delegation is recalled and the file returned.
- Compound procedures (e.g. lookup file, open file, and read data all in one call instead of 3). Note that NFS uses RPCs internally.
- Exporting different parts of a file system to different remote servers.

7.3 Google File System (GFS)

GFS is a DFS that stripes files across multiple commodity servers, and is layered on top of ordinary Linux file systems. Although GFS is proprietary, the Hadoop Distributed File System (HDFS) is a simplified open-source implementation of GFS. Below are a few key properties of GFS:

- **Synchronization:** A GFS master stores data about files and chunks. This metadata is cached in the main memory of chunk servers, and updates to these chunks servers are written to their own main memory. The master periodically polls the chunk servers to keep this metadata consistent.
- **Reads:** For reads, a client sends a file name and chunk index to the master, which responds with an address for the server storing that chunk.
- **Writes:** Updates are written directly to chunk servers, after which the changes are propagated through all primary and then secondary replicas.

7.4 File Sharing Semantics

In centralized file systems, reads and writes are strictly ordered in time. This ensures that an application can always read its own writes. This behavior is often described as UNIX semantics, and can be attained in a DFS so long as there is a single file server and files are not cached.

When a cached file in a DFS is modified, it is impractical to immediately propagate the changes back to the remote server (this would largely defeat the purpose of caching in the first place). Instead, a widely implemented rule is that changes are only visible to the process or machine that modified a file, and only made visible to other processes or machines when the file is closed. This rule is known as session semantics.

Semantics of DFS can be defined using combinations of various techniques. For example, NFSv4 supports session semantics and byte range file locking, and HDFS provides immutable files along with an append function (e.g. for storing log data).

8 Hadoop MapReduce

8.1 Background

MapReduce frameworks address the problem of parallelizing computations on big data sets across many machines. Google did much of the initial work on a generic MapReduce framework intended to phase out the need for special-purpose frameworks for different kinds of computations. Hadoop MapReduce is the most prominent open-source implementation of Google's MapReduce framework, and was created by Yahoo engineers Doug Cutting and Mike Cafarella.

8.2 High-level Architecture

Hadoop consists of a MapReduce engine and an HDFS system. The MapReduce layer contains a JobTracker, to which clients can submit MapReduce jobs. This JobTracker pushes work to available TaskTracker nodes, with the goal of picking a node close to the data that job requires.

8.3 MapReduce Basics

A few key guiding principles of MapReduce are:

- Components do not share data arbitrarily, as the communication overhead for keeping data synchronized across nodes can be very high.
- Data elements are immutable. Computation occurs by generating new outputs, which can then be forwarded into the next computation phase.
- MapReduce transforms lists of input data elements into lists of output data elements. A single MapReduce program typically does so twice, once for a Map phase and once for a Reduce phase.

8.3.1 Mappers

A list of data elements is provided to a mapper, which transforms each element to some output element.

8.3.2 Reducers

A reducer receives an iterator of input values, and combines these values together to produce a single value.

8.4 Technical Definitions

- **InputSplit:** A unit of work assigned to a single map task.
- **InputFormat:** Determines how input files are parsed, which also determines the InputSplit.
- **RecordReader:** Loads data from input split, creating key-value pairs used by the mapper.
- **Partitioner:** Determines which partition key-values pairs should go to.
- **OutputFormatter:** Determines how output files are formatted.
- **RecordWriter:** Writes records to output files.

8.5 Fault Tolerance

The main way Hadoop achieves fault tolerance is by restarting failed tasks. TaskTracker nodes are in constant communication with the JobTracker node, so should a TaskTracker fail to respond in a given period of time, the JobTracker will assume it has crashed. If the job that crashed was in a mapping phase, then other TaskTracker nodes will receive requests to re-run all map tasks previously run by the failed node. If the job that crashed was in a reduce phase, then other TaskTracker nodes will receive requests to re-run all reduce tasks that were in progress on the failed node.

Such a simple fault tolerance mechanism is possible because mappers and reducers limit their communication with one another and the outside world. A potential drawback with such an approach is that a few slow nodes (called stragglers) can create bottlenecks that hold up the rest of the program. One strategy to remedy this problem is known as speculative execution, where the same task is assigned to multiple nodes to decrease the expected time in which it will be finished.

8.6 Common Programming Patterns

8.6.1 Counting and Summing

The simplest mapper for a counting problem would output a (key, 1) tuple for an instance of a given key. Alternatively, the mapper can aggregate data for an entire document, or a combiner can be run just after the mapper to aggregate data across documents processes by a map task in a similar way to how a reducer would.

8.6.2 Selection

Selection returns a subset of input elements that meet a certain requirement. Selection can be handled entirely in the map task; a reducer need not be specified, in which case one output file will be generated per map task.

8.6.3 Projection

Projection returns a subset of fields of each input element (e.g. a, b, c becomes a, b). A mapper can handle projection on individual elements, while a reducer is only needed to eliminate duplicates.

8.6.4 Index Inversion

Index inversion produces a mapping of terms to document ids. Mappers emit (term, id) tuples, while reducers combine these tuples to generate lists of ids for each term.

8.6.5 ?Cross-Correlation

Cross-correlation problems provide as input a set of tuples of items, and for each possible pair of items, the number of items where the tuples co-occur is measured. If N items are provided, then N^2 values should be reported. This problem can be represented through a matrix if N^2 is small enough. For larger values of N^2 , MapReduce can be effective.

A pairing implementation where mappers return tuples with a pair and a 1 count is simple but often not performant. A “stripes” approach where mappers return tuples with an item and a list of items it appears with tends to perform better despite requiring more memory for map-side aggregation, since there end up being roughly N keys (one for each element) instead of N^2 (one for each pair).

9 Apache Spark

9.1 Background

Cluster computing frameworks (e.g. Hadoop) allow for large-scale data computations that automatically handle work distribution and fault tolerance. However, many of these frameworks do not leverage distributed memory efficiently, making them ineffective for computations with intermediate results. Apache Spark introduces the concept of resilient distributed datasets (RDDs), which are parallel, fault-tolerant data structures that persist intermediary results in memory.

9.2 Lineage

A lineage is a model of the flow of data between RDDs. Spark designs a lineage to perform efficient computations, and also uses the lineage to determine which RDDs to rebuild to recover from failures.

9.3 Transformations and Actions

Transformations are data operations that convert an RDD or a pair of RDDs into another RDD. Actions are data operations that convert an RDD into an output. When an action is invoked on an RDD, the Spark scheduler puts together a DAG of transformations.

9.3.1 Narrow vs Wide Dependencies

The transformations in the DAG are grouped into stages. A single stage is a collection of transformations with narrow dependencies, meaning that one partition of the output depends only on a single partition of the input (e.g. map, filter). The boundaries between stages feature wide dependencies, meaning that a single partition of output may correspond to multiple partitions of input (e.g. group by key), so the transformations may require a shuffle.

9.3.2 Pagerank

???

9.4 Spark vs Hadoop MapReduce

Two common differences between Spark and MapReduce are:

- Spark stores intermediary results in memory, while MapReduce dumps results to HDFS between each job. The MapReduce approach can lead to unnecessary I/O when running multiple jobs in sequence.
- Spark can support some more complicated workflows within a single job rather than running several new ones.

10 Distributed Graph Processing

Many data sets (e.g. web hyperlinks, social networks, transportation routes) can be modeled using graphs. Computations concerning very large graphs can benefit from distributed architectures. However, previously examined cluster computing frameworks (e.g. Hadoop MapReduce, Apache Spark) are typically not good fits for graph processing, which motivates the creation of separate distributed graph processing frameworks.

10.1 Google’s Solution: Pregel

Pregel is a proprietary framework developed by Google to perform computations on large distributed graphs with performance and ease of programming in mind. A master/worker model similar to those of MapReduce and Spark is used, where each worker is responsible for a particular vertex partition. The framework maintains some state information for each vertex, including:

- Problem-specific values
- List of messages sent to vertex
- List of outgoing edges from a vertex
- A binary active/inactive state

10.1.1 Supersteps

Pregel applies a bulk synchronous parallel (BSP) model of computation that where a computation is divided into iterative rounds called supersteps. Workers perform their computations asynchronously within each superstep, and only exchange data between supersteps. Specific actions that can be taken by a worker within a superstep include:

- Receiving messages sent in the previous superstep.
- Sending messages to be received in the next superstep.
- Modifying vertex values or edges.
- Deactivating a vertex, which is reactivated when it receives a message.

This distributed execution stops when all vertices are inactive and with no more messages to process.

10.1.2 Initialization

The master is responsible for assigning a section of the vertices to each worker. The default partitioner uses a simple hash function over vertices, which generates a fairly even distribu-

tion of vertices. To take advantage of other properties of a graph (e.g. exploiting locality), the default partitioning scheme can be overridden.

10.1.3 Combiners and Aggregators

Pregel supports combiners, which serve to reduce the amount of data exchanges over a network and the number of messages. Combiners are often applicable when the function applied at each vertex is commutative and associative (e.g. sum, max).

Pregel also supports aggregators, which are used to compute aggregate statistics from vertex-reported values.

10.1.4 Fault Tolerance

At the end of each superstep, the master instructs workers to save their state (i.e. vertex values, edge values, incoming messages, aggregator values, etc.) to persistent storage. When the master detects the failure of a worker node, it rolls back all workers to the last successful superstep, and the computations resume. If deterministic replay of sent messages is possible, then there are some more efficient recovery mechanisms that only require that the failed worker is rolled back to the last successful superstep.

10.2 Open-source Implementations

Pregel-like APIs are supported by various open-source frameworks such as Apache Giraph. Other APIs support centralized graph processing (e.g. GraphChi), which tends to perform better for fairly small datasets that can fit in the main memory of a single machine.

11 Consistency and Replication

11.1 Replication Overview

11.1.1 Purposes of Replication

The main purposes behind data replication include:

- **Increasing reliability**, since data is more likely to be accessible if it exists on more replicas.
- **Increasing throughput**, since replicas can serve read operations in parallel.
- **Lowering latency**, since replicas physically close to clients can avoid costly round trips.

11.1.2 Replicated Data Stores

In replicated data stores, each data object is replicated across multiple hosts. Replicas may reside on the same host as processes that interact with them, or may be remote.

11.2 Consistency Models

Consistency models describe the extent to which replicas can disagree on the current state of data. Managing consistency tends to be straightforward for read-only data, but the need for sophisticated models arises once replicas hold mutable states, especially ones shared among multiple processes. Selecting a good consistency model is often challenging, as application requirements rarely map neatly to a specific model.

11.2.1 Sequential Consistency

Sequential consistency requires that:

1. The result of any execution is the same as if the operations by all processes on the data store were performed in some sequential order.
2. The operations of each individual process appear in the order specified in its program.

The notion of this consistency model is borrowed from shared memory multiprocessing.

11.2.2 Causal Consistency

The concept of causal consistency is also borrowed from shared memory multiprocessing, and is based on a notion of similarity similar to Lamport's "happens before" model. A data

store is causally consistent if writes related by some “causally precedes” operation are seen by all processes in the same order. The “causally precedes” operator is defined as follows:

1. Op1 causally precedes Op2 if Op1 occurs before Op2 in the same process.
2. Op1 causally precedes Op2 if Op2 reads a value written by Op1.

11.2.3 Linearizability

A data store is linearizable when the result of any execution is the same as if all operations on the data store were executed in some sequential order that extends Lamport’s “happens before” model. This means that if Op1 finishes before Op2 begins, then Op1 must precede Op2 in sequential order.

11.2.4 Eventual Consistency

The idea behind eventual consistency is that if no updates take place for some time, all replicas will gradually become consistent. Eventual consistency allows for different processes to observe write operations taking effect in different orders, even when the operations are related by “causally precedes” or “happens before”.

Eventual consistency is a very weak property on its own, and appears to guarantee very little if updates are applied continuously. To strengthen these guarantees, session guarantees can be applied, which restrict the behavior of a single process in a single session. Examples of session guarantees include:

- **Monotonic reads**, which means that if a process reads a particular value of a data item, then successive reads in that process will return that value or a newer one.
- **Read your own writes**, which means the effect of a write operation will always be seen in successive read operations by the same process.

11.3 Replication Protocols

11.3.1 Primary-Based Replication Protocols

Primary-based replication protocols involve updates being performed on a single primary replica, with updates then pushed to any backup replicas. If a primary replica is stationary, where it may be often updated remotely by other servers, this protocol can be classified as remote-write. Alternatively, local updates may be enabled by allowing the primary replica to migrate from server to server, in which case the protocol can be classified as local-write.

Primary-based replication protocols allow for some strong consistency models. However, a single primary replica can lead to performance bottlenecks and availability problems.

11.3.2 Quorum-Based Replication Protocols

Quorum-based replication protocols allow all replicas to perform updates and to service reads. However, each update and read must be performed on a sufficiently large subset of replicas, known as the write and read quorum respectively.

Let N be the total number of replicas, N_R be the read quorum, and N_W be the write quorum. In a distributed database, the following two rules must be satisfied:

1. $N_R + N_W > N$, which enables the detection of read-write conflicts.
2. $N_W + N_W > N$, which enables the detection of write-write conflicts.

Subsets of replicas that do not satisfy the rules above are referred to as partial quorums. Alternate means must be found for resolving conflicts, such as tagging updates with timestamps.

11.3.3 Eventually-Consistent Replication Protocols

In simple cases of eventual consistency, reads and updates are serviced by the closest replica. This replica lazily (i.e. async to original update) propagates updates to other replicas. Replicas can go out of sync for various reasons, such as one being unreachable due to network failure. This is handled by what is known as an anti-entropy mechanism. One sample anti-entropy mechanism is different replicas exchanging hash trees (also known as Merkle trees), which contain hashes of blocks of data that can be used to find where two data sets differ.

12 Fault Tolerance

12.1 Dependability

Fault tolerance is closely related to the concept of dependability, which implies the following requirements:

- **Availability:** The system should operate correctly at a given instant in time (e.g. 99% availability means a system operates correctly 99% of the time).
- **Reliability:** The system should run continuously without interruption (e.g. mean time between failures (MTBF) of one month).
- **Safety:** System failures should not have catastrophic consequences.
- **Maintainability:** A failed system should be easy to repair.

12.2 Failures

A failure is occurring when a system cannot fulfill its promises. Five major types of failures of servers in distributed systems are:

- **Crash failure:** A server is working correctly and then halts.
- **Omission failure:** A server fails to respond to requests or fails to receive or send messages.
- **Timing failure:** A server's response falls outside of some specified time window.
- **Response failure:** A server's response is incorrect.
- **Arbitrary failure:** A server produces arbitrary failures at arbitrary times.

12.3 Errors and Faults

An error is a part of a system's state that may lead to failure (e.g. corrupt data read from hard disk). The fault is the underlying cause of the error (e.g. hard disk head crashes). The main types of faults are:

- **Transient faults**, which emerge once and then disappear (e.g. birds flying in front of a microwave receiver)
- **Intermittent faults**, which tend to reappear (e.g. loose contact on electrical connector)
- **Permanent faults**, which continue to exist until some faulty component is replaced (e.g. burned out power supply)

12.4 Masking Failure by Redundancy

Triple-modular redundancy (or more generally, N-modular redundancy) is an approach to fault tolerance where multiple (3 or N) subsystems perform the same process, after which the result is determined by a majority voting system. Thus, the system can continue to function if some of its subsystems fail so long as the remaining systems can win a majority vote. To support voting, processes must be able to communicate with one another. This communication can be setup as a flat group (i.e. complete graph between all processes) or in some sort of hierarchical structure with a coordinator process and multiple worker processes.

12.5 Consensus Problem

The consensus problem in distributed computing can be defined as follows:

- Each process has procedures `propose(val)` and `decide()`
- Each process first proposes a value by calling `propose(val)` once
- Each process learns the value agreed upon by calling `decide()`

The following properties must hold:

- **Agreement:** Two calls to `decide()` cannot return different values.
- **Validity:** If `decide()` returns a value `val`, then some process called `propose(val)`.
- **Liveness:** Calls to `decide()` and `propose(val)`, if a process does not fail, must eventually terminate.

12.5.1 Variations

The solvability of the consensus problem depends on several factors of a distributed environment, including:

- Whether processes are synchronous or asynchronous to one another
- Communication delays between processes
- Message delivery order
- Whether messages are unicast (one-to-one) or multicast (one-to-many)

12.6 RPC Semantics Under Failures

RPC systems may exhibit the following failure scenarios, which depend on when a failure occurs in an RPC call:

- Client unable to locate server.

- Request message from client to server is lost.
- Server crashes after receiving request.
- Reply message from server to client is lost.
- Client crashes after sending request.

12.6.1 Dealing with RPC Server Crashes

When the server crashes after receiving a request, the client may not know whether the server crashed before or after executing the request. Techniques for dealing with such RPC server crashes include:

- Reissuing the request, leading to at-least-once semantics, which are safe as long as the request is idempotent (i.e. repeated executions have the same effect as one execution).
- Not reissuing the request and reporting a failure, leading to at-most-once semantics, with no guarantee that the request was processed.
- Determining whether the request was processed, and if it was not, reissuing the request. This leads to exactly-once semantics, which may be difficult to implement as there may be no clear way of knowing what action was performed.

13 Apache ZooKeeper

ZooKeeper is a centralized system that manages distributed systems as a hierarchical key-value store. Common uses of ZooKeeper include maintaining configuration information, providing distributed synchronization, and managing naming registries for distributed systems. ZooKeeper emphasizes good performance (particularly for read-dominant workloads), being general enough to be used for many different purposes, reliability, and ease of use.

13.1 Data Model

ZooKeeper's data model is a hierarchical key-value store similar to a file system. Nodes in this store are called *znodes*, and may contain data and children. Reads and writes to a single node are considered to be atomic, with values read or written fully or not at all.

13.2 Node Flags

Two important flags nodes can carry are:

- **Ephemeral flags**, which make nodes exist as long as the session that created them is active, unless they were explicitly deleted.
- **Sequence flags**, which make nodes append a monotonically increasing counter to the end of their path.

13.3 Consistency Model

ZooKeeper ensures that writes are linearizable and that reads are serializable (a similar property to sequential consistency). ZooKeeper also guarantees per-client FIFO servicing of requests.

13.4 Servers

When running in replicated mode, all servers have a copy of the state in memory. A leader is elected at startup, and all updates go through this leader. Update responses are sent once a majority of servers have persisted the change. Thus, in order to tolerate n failures, $2n + 1$ replicated servers are required.

14 Distributed Commits and Checkpoints

14.1 Transactions

Transactions are indivisible operations with the following properties:

- **Atomic:** An operation occurs fully or not at all. This can be hard to achieve in distributed environments.
- **Consistent:** A transaction is a valid transformation of the state.
- **Isolated:** A transaction is not aware of other concurrent transactions.
- **Durable:** Once a transaction completes, its updates persist, even in the event of failure.

14.2 Two-Phase Commits

The two-phase commit (2PC) is a distributed transaction commitment protocol. Systems running this protocol consist of a coordinator and participants.

In the first phase, the coordinator asks participants whether they are ready to commit. These participants respond by voting. In the second phase, the coordinator analyzes the participant votes. If all participants voted to commit, the commit proceeds. Otherwise, the transaction is aborted.

The 2PC procedure described above makes a few key assumptions, including that processes are synchronous and that communication delays are bounded.

14.2.1 Recovery from Failures

Participants are able to make progress so long as they have received a decision. This decision is typically received from the coordinator, though in the case of a coordinator crash, may be received from another participant.

The protocol described above blocks when all participants are waiting for an answer and the coordinator crashes, though some more complicated implementations can overcome this. However, even in these more complicated implementations, the simultaneous failure of the coordinator and some participants can make it difficult to detect whether all participants are ready for an answer.

14.2.2 Distributed Checkpoints

Recovery after failure is only possible if the collection of checkpoints by individual processes forms what is called a distributed snapshot. A distributed snapshot requires that process

checkpoints (i.e. representations of state at a certain point) contain a corresponding send event for each message received. The most recent distributed snapshot is called the recovery line.

Coordinated checkpointing algorithms can be applied to create recovery lines. A sample two phase checkpointing algorithm works as follows:

1. The coordinator sends a checkpoint request message to all participant processes. When a participant receives the message, it pauses processing of incoming messages, forms a checkpoint, and returns an acknowledgment (ACK) to the coordinator.
2. Once all ACKs are received by the coordinator, it sends a message to the participants that they can resume processing incoming messages.

15 RAFT

16 Apache Kafka

Apache Kafka is an open-source distributed stream-processing platform. A few key features of Kafka are publish-subscribe message oriented communication, real-time stream processing and distributed and replicated storage of these messages and streams. Common use cases for Kafka include:

- High-throughput messaging
- Activity tracking
- Metric collection
- Log aggregation

16.1 Topics and Logs

A topic is a stream of records. Each topic can have multiple partitions, which allow a topic to be parallelized by splitting its data among multiple Kafka brokers (i.e. servers), which in turn can have multiple replicas. Each record in a partition has a unique offset.

16.2 Producers and Consumers

Producers push records to Kafka brokers for a specific partition. These pushes can be batched (i.e. multiple records delivered in one operation), and idempotent delivery is supported (i.e. duplicate commits can be prevented).

Consumers pull batches of records from Kafka brokers. Each consumer has a position within a partition, which can be updated once the records are processed.

16.3 Stream Semantics

The main two ways of interpreting streams are:

- **Record streams**, where each record represents a single event (e.g. a single purchase).
- **Changelog streams**, where each record represents an update to a state (e.g. a bank account balance change).

Kafka provides different APIs for interacting with the different types of streams. The KStream API is used to interact with record streams, while the KTable API is used to interact with changelog streams. The KGroupedStream API can be used when converting from a KStream to a KTable.

16.4 Windowed Streams

Windowing allows for grouping records close in time. A few common types of windows include:

- **Hopping time windows**, which are defined by a size and advance interval (hop). An example is to compute something every 10 seconds (hop) for the last 60 seconds of data (size).
- **Tumbling time windows**, which are special non-overlapping versions of hopping time windows (i.e. the hop and size are equal).
- **Sliding windows**, which slide continuously over the time axis.
- **Session windows**, which aggregate data by activity periods (i.e. new windows are created once a period of inactivity exceeds a certain length).

17 Clock Synchronization

Clocks are instruments used to measure time. A lack of synchronization among clocks in a distributed system can result in unclear ordering of events.

17.1 Background on Clocks

17.1.1 Calendars

A few significant calendars used historically include:

- **Roman Calendar:** The Roman calendar is a lunar calendar (i.e. based on moon phases) with months of 29 or 30 days. The original Roman calendar had only 10 months, with winter days unallocated. Reforms such as adding two extra months were later introduced, but the calendar remained difficult to align with seasons, since it is a lunar calendar.
- **Julian Calendar:** The Julian calendar was the first solar calendar (i.e. based on the Earth's movement around the sun). Introduced in 45 BC, the calendar improved on the Roman calendar, such as by adding the concept of a leap year (but not very accurately).
- **Gregorian Calendar:** The Gregorian calendar is the most widely used calendar in the world. Introduced in 1582, its main improvement over the Julian calendar is measuring leap years far more accurately (accurate to 1 day in 7700 years).

17.1.2 Timekeeping Standards

A few standards used for timekeeping include:

- **Solar day:** The interval between two consecutive transits of the sun (i.e. sun reaching highest point of the day). This is not constant.
- **TAI (Temps Atomique International):** International time scale based on the average of multiple Cesium 133 atomic clocks.
- **UTC (Coordinated Universal Time):** The world's primary time standard. Based on TAI, UTC uses leap seconds at irregular intervals to compensate for the Earth's slowing rotation. As a result, UTC is currently tens of seconds behind TAI.

17.1.3 Limitations of Atomic Clocks

Atomic clocks are the most accurate timekeeping devices known, but are limited by relativistic time dilation. This was confirmed by the Hafele-Keating experiment, in which atomic clocks were flown around the world and compared to stationary ones.

17.2 Terminology

Let $C(t)$ denote the value of a clock C at a reference time (e.g. UTC) t .

- **Clock skew** of C relative to t is $\frac{dC}{dt} - 1$.
- **Offset** of C relative to t is $C(t) - t$.
- **Maximum drift rate** of C is a constant ρ such that $1 - \rho \leq \frac{dC}{dt} \leq 1 + \rho$.

17.3 Network Time Protocol (NTP)

NTP is a common time synchronization protocol for computers over variable-latency data networks.

Consider two hosts (A and B). Host A sends a request at T_1 to host B. Host B receives this request at T_2 . Host B replies at T_3 , and host A receives this reply at T_4 . The offset of B relative to A is $\theta = \frac{(T_2 - T_1) + (T_3 - T_4)}{2}$. The estimated one-way network delay is $\delta = \frac{(T_4 - T_1) - (T_3 - T_2)}{2}$. NTP typically collects multiple (θ, δ) pairs and uses the minimum value of δ as the best estimation of the delay. The corresponding θ is considered to be the most reliable estimate of the offset.

17.3.1 Clock Strata

A reference clock, such as an atomic clock, is said to operate at stratum 0. A computer at stratum 1 has its time controlled by a stratum 0 device, a computer at stratum 2 has its time controlled by a stratum 1 device, and so forth. Computers may only adjust the time of computers at greater strata, and when adjusting, must make sure time does not appear to flow backwards.

17.3.2 Precision

NTP accuracy is typically measured in 10s of milliseconds. Some other protocols, such as Precision Time Protocol (PTP), can be used to achieve sub-millisecond accuracy where NTP does not suffice.

17.4 Logical Clocks

Logical clocks are not concerned with the precise timing of events, but instead make sure that the time at each process follows a “happens before” relation (e.g. a message cannot be received before it is sent).

17.4.1 Lamport Timestamping

Lamport timestamping is a technique to enforce the “happens before” relation described above. The algorithm is:

- Before executing an event (send or receive), a process increments a counter.
- Sent messages are tagged with the sender’s timestamp.
- Upon receiving a message, a process sets its counter to be the max of its current counter and the received counter (and increments the result).

17.4.2 Vector Clocks

Vector clocks represent logical time similar to Lamport timestamping. However, instead of a single counter, each clock stores a vector of a counter at each process. Vector changes are similar to those of Lamport timestamping, but applied to a single entry in the vector instead of a process-wide counter.

Lamport timestamping guarantees that if an event a happened before b , then the clock at a process that handled both a and b is smaller for a . However, the reverse of the relationship does not hold for Lamport timestamping (i.e. a lesser time at a than b does not imply that a happened before b). Vector clocks guarantee the reverse as well, so they are said to provide a complete characterization of causality.

18 CAP Principle

The CAP theorem, also known as Brewer’s theorem, states that it is impossible for a distributed system to provide all three of:

- **Consistency**, which means that nodes in a system agree on the most recent state of data.
- **Availability**, which means that nodes are able to execute read-only queries and updates.
- **Partition tolerance**, which means that the system continues to function if its servers are separated into disjoint sets (e.g. because of a network failure).

18.1 CP vs AP Systems

CAP can be interpreted as that when a network partition occurs, systems need to choose between consistency (C) and availability (A). CP systems select consistency over availability, while AP systems select availability over consistency. AP systems are appropriate for many applications that are latency-sensitive, inconsistency-tolerant, and do not rely on transactions.

CP systems are able to enforce stronger consistency models (e.g. serializability, linearizability) than AP systems, which may strive for eventual or causal consistency.

18.2 PACELC

PACELC is an extension of the CAP theorem that states that:

- If there is a **P**artition, a choice must be made between **A**vailability and **C**onsistency.
- **E**lse, a choice must be made between (lower) **L**atency and **C**onsistency.

18.3 Quorum-Based Consistency and CAP

Systems that ensures every read observes the effects of all writes that completed before the read started, such as through using quorums, are said to have strong consistency. This type of consistency is considered to be “C” in the context of the CAP theorem.

Some systems allow consistency settings to be determined on a per-request basis (e.g. one, quorum, all). These systems are said to have tunable consistency.

In general, systems where a read or write must occur at specific machines are not available. One solution to this is using sloppy quorums, in which a set of replicas can change dynamically, such as to adjust to network partitions. This is used in database systems with hinted

handoff, where an arbitrary node can accept an update; if that update should ideally go to a different node, the update will be transferred to that node once the network partition is removed.